

CSE 444: SimpleDB Final Report

Linxing Preston Jiang Winter 2018

March 12, 2018

1 Overall System Architecture

In a typical database architecture, there are four main components: Process Manager, Query Executor, Share Utilities, and Storage Manger (Balazinska, Mass, Lecture 3). For our implementation of SimpleDB in the labs, we focus on the Storage Manager and Query Executor: adding access methods for data store on disk in lab1, adding both file mutability (insert/delete tuples to/from file system, eviction from full BufferPool) and query operators (SeqScan, Join, etc.) & aggregates (min, max, etc.) in lab2, adding lock manager in lab3, adding log manager in lab4, and finally, adding parallel data processing in lab6.

Figure 1 shows the parts of the architecture of SimpleDB which we implemented in the labs.

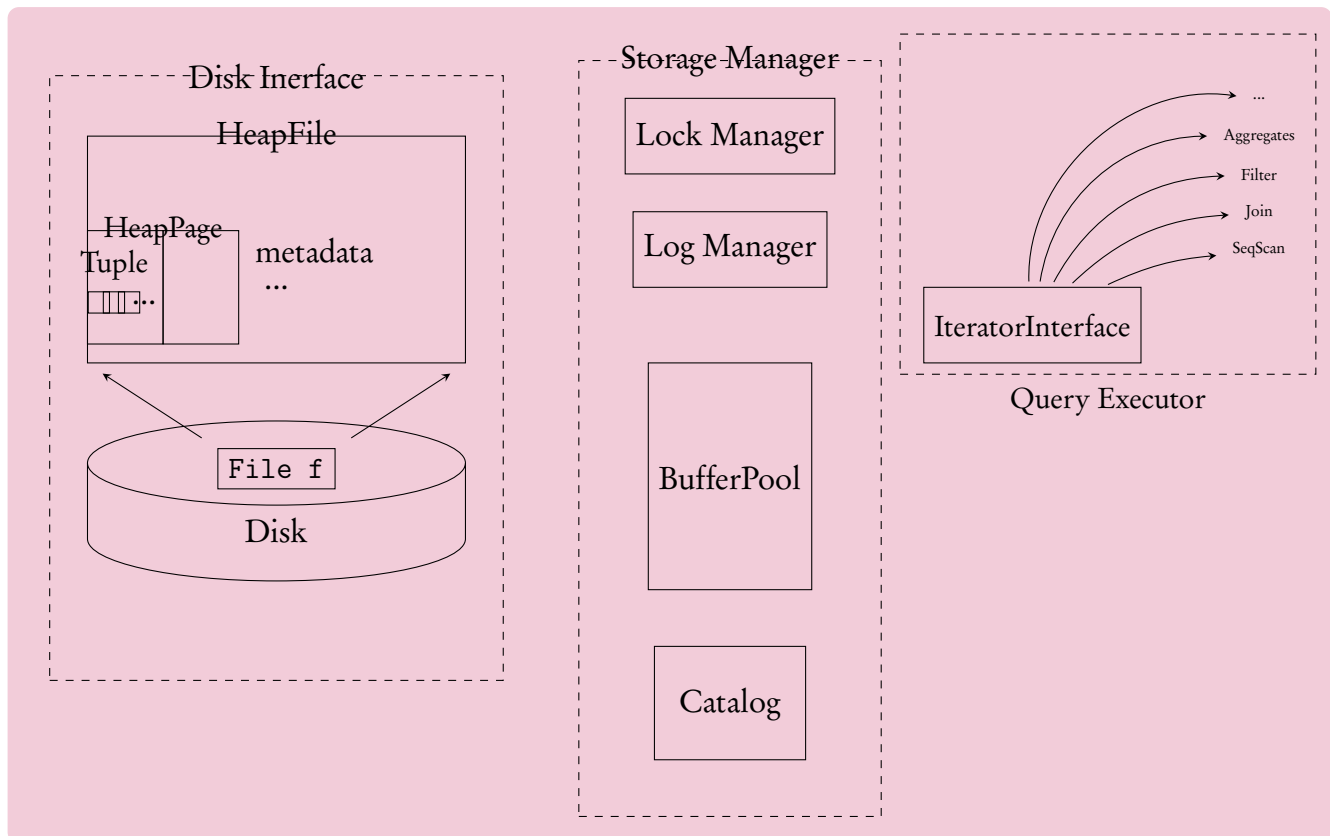


Figure 1: SimpleDB Architecture

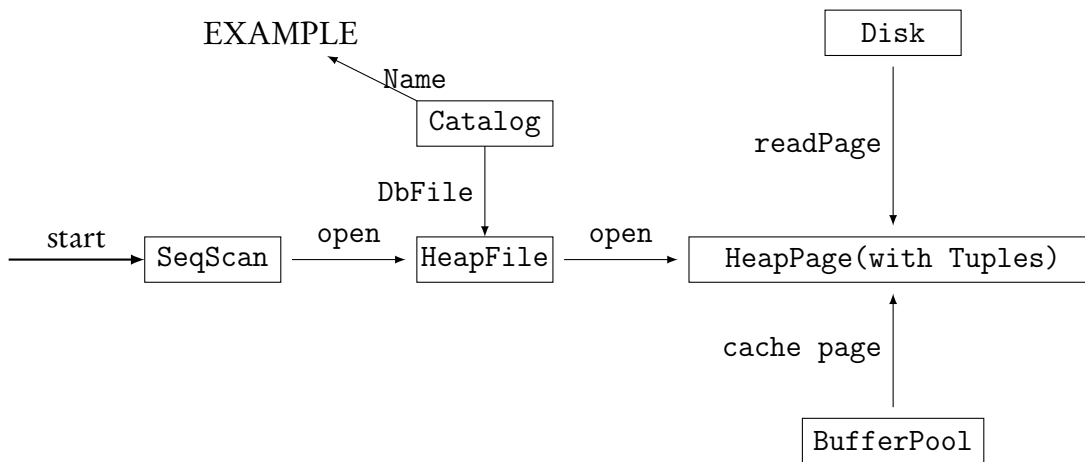


Figure 2: SimpleDB open

1.1 BufferPool and Operators

BufferPool is responsible for both caching pages in memory that have been recently read from disk and handle concurrency and transactions. All operators read and write pages from various files on disk through the buffer pool. Operators are responsible for executing query plans. In SimpleDB, each operator implements the `OpIterator` interface, which supports `open`, `hasNext`, `next`, `rewind`, and `close`. Operators are connected together into a plan by passing lower-level operators into the constructors of higher-level operators (Lab1 ReadMe). Programs call `next` on the root operator and then fetch tuples recursively through the plan tree in one pass top-down and another pass bottom-up.

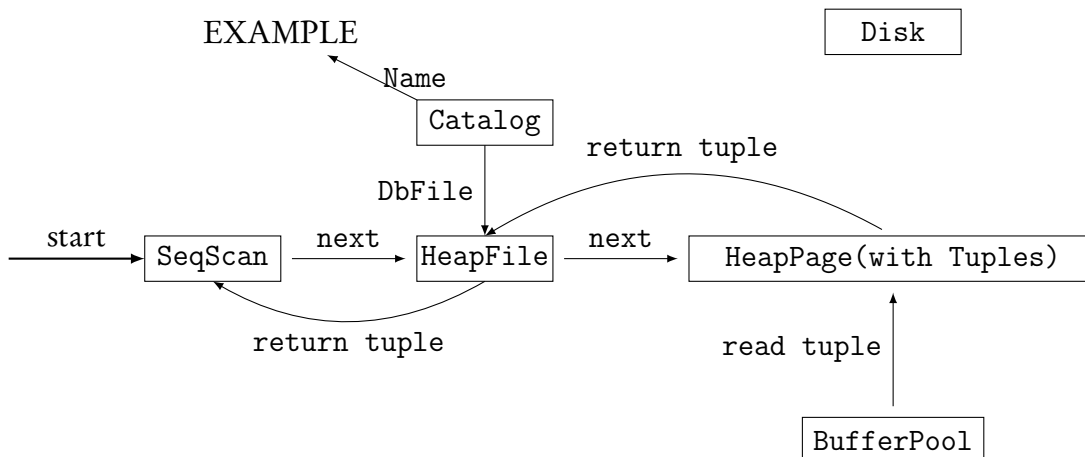


Figure 3: SimpleDB next

In Lab 1, we implemented `getPage` which is needed for reading Pages into memory and getting tuples, together with `SeqScan` operator to scan the file and return tuples. The workflow of opening operators to read file and caching pages into BufferPool is shown by Figure 2, and the workflow of recursively getting tuples through `SeqScan` using `next()` is shown by Figure 3.

In Lab 2, We added insert, delete functionalities in SimpleDB, which insert/delete tuples to/from the

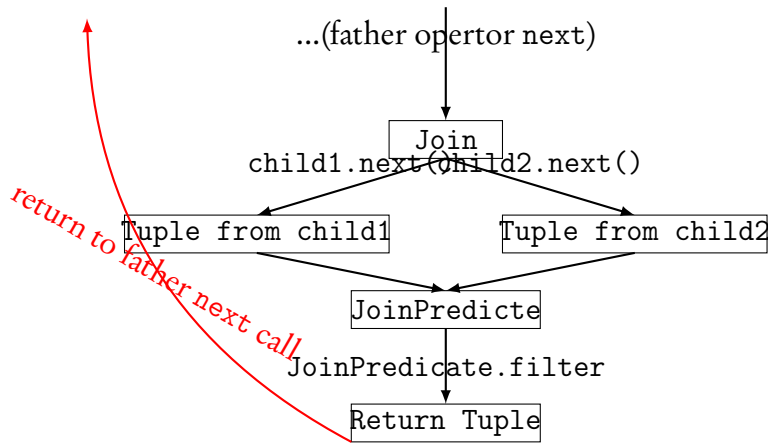


Figure 4: SimpleDB join

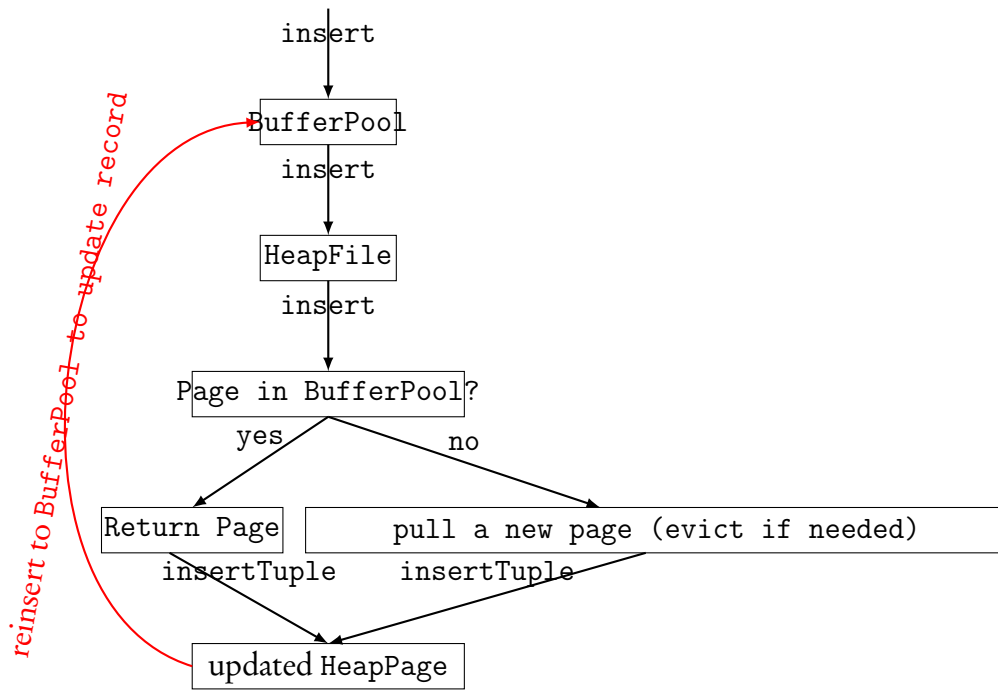


Figure 5: SimpleDB insertTuple

pages in BufferPool by calling the method from HeapFile to update the page, then BufferPool updates the records by re-inserting the pages into the BufferPool. When the BufferPool is full and a new page need adding, writePage from HeapFile will be called to write the dirty page to disk and add the new page into the BufferPool. BufferPool is in charge of updating the pages because Insert/Delete operators directly call insert/delete of BufferPool. Besides, we also implemented Filter, Join operators and the aggregates. As an example, Figure 4 shows the workflow of using Join operator and Figure 5 shows the workflow of inserting tuples.

2 Parallel Data Processing

3 Discussion