

CSE 444: Homework 5

Linxing Preston Jiang Winter 2018

February 7, 2018

1 Query Plan Cost Computation

1. (a) $\frac{1000-100}{9000-0} = \frac{1}{10}$
(b) $\frac{1}{V(S,e)} \cdot \frac{1}{V(S,f)} = \frac{1}{10} \cdot \frac{1}{100} = \frac{1}{1000}$
(c) $\frac{1}{\max\{V(R,c), V(S,d)\}} = \frac{1}{50}$
(d) $\frac{1}{V(R,b)} + \frac{1}{V(R,b)} = \frac{1}{50}$
(e) $\frac{1}{\max\{V(S,g), V(T,h)\}} = \frac{1}{100}$
2.
 - $|R_1| = 10000 \cdot 0.1 = 1000$
 - $|R_2| = 10000 \cdot \frac{1}{1000} = 10$
 - $|R_3| = \frac{1000 \cdot 10}{50} = 200$
 - $|R_4| = \frac{200}{50} = 4$
 - $|R_5| = \frac{4 \cdot 10000}{100} = 400$
 - $|R_6| = |R_5| = 400$
3. (a) Because of the clustered index on a , we have $B(R) \cdot 0.10 = 100$ IOs to read from R . And the unclustered index on (e, f) will help reduce the number of IOs to $B(S) \cdot \frac{1}{1000} = 10$ pages. We do not write intermediate results to disk, so the next IO cost comes from reading from T , which will have $T(R_4) \cdot \frac{B(T)}{V(T,b)} = 4 \cdot \frac{1000}{100} = 40$. There the total cost is

$$100 + 10 + 40 = 150$$

2 Query Optimization


1. See table

Subquery	Cost	Size of output	Plan	P/K
R	100 page IOs	1K records on 100 pages	Sequential scan of R	K
S	1K page IOs	10K records on 1K pages	Sequential scan of S	K
W	10 page IOs	100 records on 10 pages	Sequential scan of W	K
RS	$100 + 100 * 1000 = 100100$ page IOs	10K records on 2K pages	Nested loop join of R (outer) S	K
RW	$100 + 100 * 10 = 1100$ page IOs	100 records on 20 pages	Nested loop join of R (outer) W	P (WR has lower IO cost)
SR	$1000 + 1000 * 100 = 101000$ page IOs	10K records on 2K pages	Nested loop join of S (outer) R	P (RS has lower IO cost)
WR	$10 + 10 * 100 = 1010$ page IOs	100 records on 20 pages	Nested loop join of W (outer) R	K
SW	P
WS	P
RSW	$100100 + 2K * 10 = 120,100$	1K records on 300 pages	Nested loop join of RS (outer) W	P
WRS	$1010 + 20 * 1K = 21,010$	1K records on 300 pages	Nested loop join of WR (outer) S	K

Cartesian Product

Best

The physical plan is shown here:



- Clustered index on R.b: Because of the selection R.b > 100, this index will help reduce the number of reads from R by the selectivity of R.b > 100.
 - Unclustered index on S.d: It depends on the estimated number of tuples from R to be read. Because we are joining R.a and S.d, if the number of tuples from R is more than the number of pages from S, we should just perform a file scan because every read from an unclustered index can be on a new page. Otherwise, we can choose to use this unclustered index to reduce the number of reads from S.
- An interesting order means an intermediate join which has a higher cost leads to a lower cost in the end. It can either be the order of tuples satisfying the order required by ORDER BY or GROUP BY, or the order on a equality-join will enable cheaper sort-merge join in the next step. For example, assume we have a hash index on R.a, a B+ tree clustered on W.c, see the following query:

```

1  SELECT *
2  FROM R, S, W
3  WHERE R.a = S.b
4         AND S.b = W.c;

```

For the step of joining R and S, an interesting order can be obtained by using a sorted-merge join, because for the next step of joining with W.c, the tuples of S.b is already sorted and so is W.c using the B+ tree clustered index. Here's another example:

```

1  SELECT SUM(R.a)
2  FROM R, S, W
3  WHERE R.a = S.b
4         AND R.c = W.c
5  GROUP BY S.b;

```

Now an interesting order is using sorted-merge-join on R.a and S.b because it will dramatically speed up the GROUP BY clause since every tuple is sorted on the GROUP BY field S.b.