

# CSE 444: Lab 1 Writeup

Linxing Preston Jiang Autumn 2017

January 17, 2018

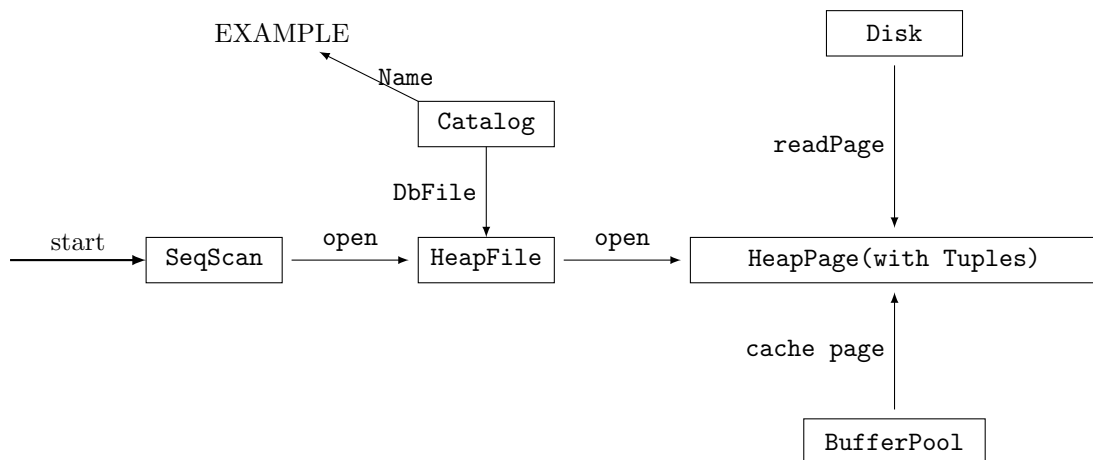
1. Lab 1 contains necessary methods for constructing databases and operators for query plan execution (just reading from file in Lab 1 with **SeqScan**). Key components implemented are:

- **Tuple**: this presents a single record (row) in a relation
- **TupleDesc**: this defines the schema of a relation (the number of fields, )
- **Catalog**: this stores information of all relations (tables) currently in the database, including the file, name, and the name of the primary key
- **BufferPool**: this caches the pages recently read by **DbFile** for good locality performance
- **HeapPage**: this stores the actual data (byte array) stored in a page
- **HeapFile**: this represents a file stored on disk as a collection of **HeapPage**
- **SeqScan**: this implements sequential scan operator for reading from tuples from **HeapFile**

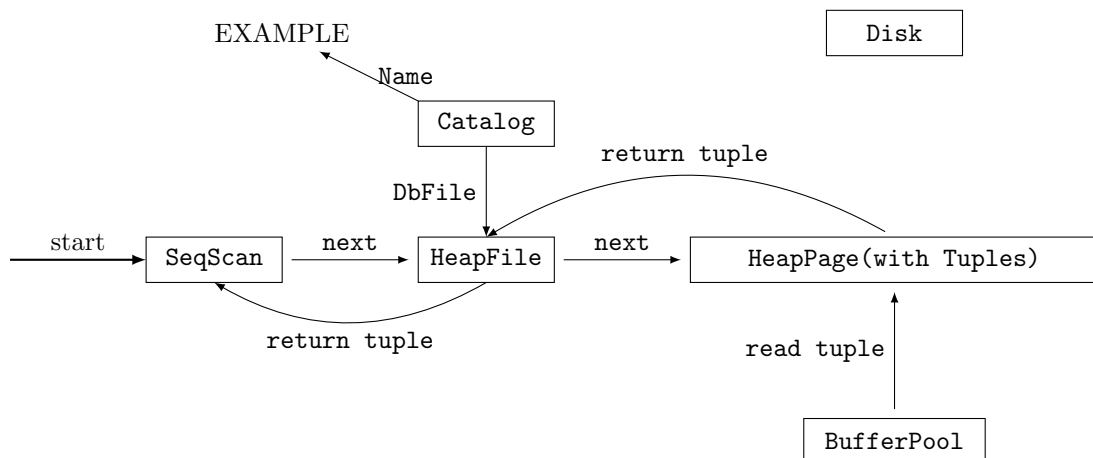
Putting them together, the workflow of a simple reading query is

SQL Query: **SELECT \* FROM EXAMPLE**

Workflow (**open()**):



Workflow (**next()**):



2. Design decisions: I followed the instructions in README and didn't make extra design decisions.
3. Extra Unit tests: A unit test for the correctness of `HeapPage.iterator()` would be nice. My initial implementation was wrong because I include the tuples in the empty slot, but the unit tests in `HeapPageIdTest` and `HeapPageReadTest` did not catch this error. It only showed up when I run `ScanTest` after finishing everything, and it was a little hard to connect the dots back to this method.
4. Changes to API: I did not make changes to the APIs given.
5. Missing elements: I believe I finished the entire Lab 1