

CSE 444: Lab 2 Writeup

Linxing Preston Jiang Winter 2018

February 4, 2018

1. Query runtime

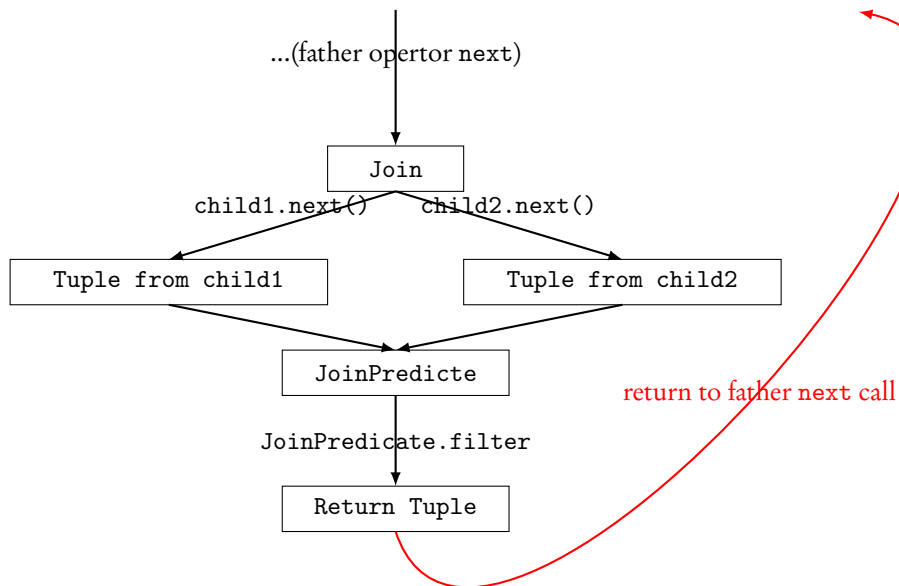
- Query 1: 1.53 seconds (attu4)
- Query 2: Because of nested-loop-join implementation, Query 2 takes forever to finish (did not finish after an eight hour tmux session on attu4)
- Query 3: Same as above

2. Lab 2 consists two major parts: operators & aggregates, and database mutability. For the first part, we implemented Filter, Join operator and common aggregates such as MIN, MAX, AVG. For the second part, we focused on the implementation of inserting/deleting tuples to and from the database, also the eviction of pages in BufferPool when it's full. Key components implemented are:

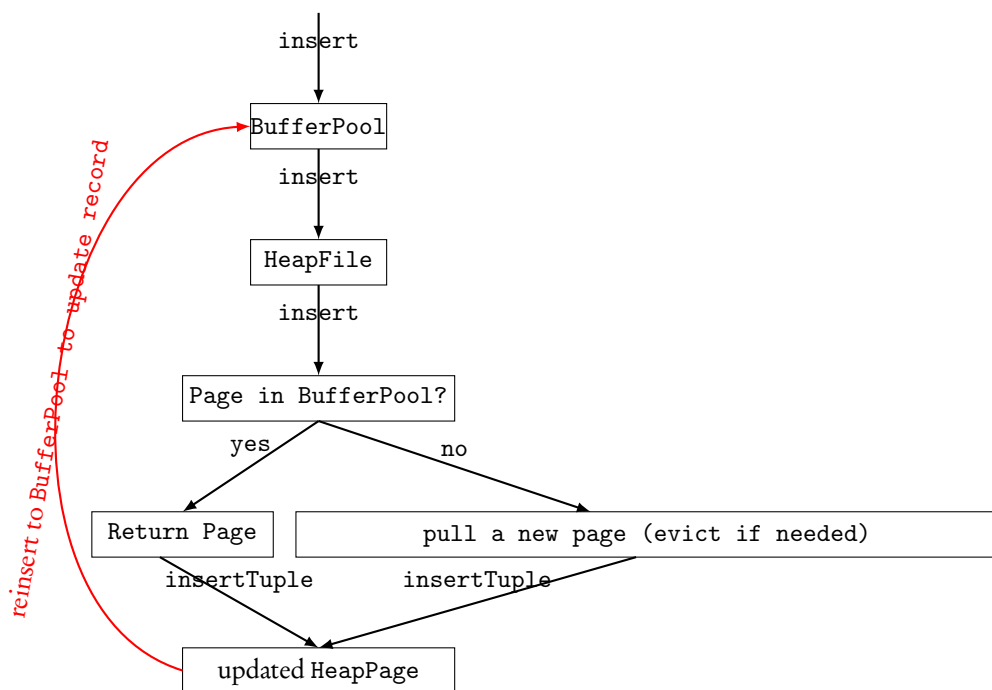
- Predicate: this defines a predicate (comparision) which is needed for operators such as Filter and Join for tuple and field comparisons. JoinPredicate is designed directly for Join operation.
- Join: The class represents the join operation for relational database. The Join operator takes in tuples from two child iterator and use JoinPredicate to filter results.
- Aggregate: this represents an aggregation clause, can be on Integer or String fields. Aggregate loads all tuples from child at open() and returns the aggregate results one group at a time in next().
- HeapPage: We added insert, delete in it, which insert/delete tuples to/from the tuple array and mark the corresponding slots to used/free. HeapPage directly inserts/deletes tuples to/from the Tuple array which will be the data updated on disk.
- HeapFile: We added insert, delete in it, which obtains the page from BufferPool (Or add a new empty page when added) and calls the insert/delete from HeapPage. Also, we added writePage method which is needed later for BufferPool to write the data on dirty HeapPages to disk.
- BufferPool: We added insert, delete in it, which insert/delete tuples to/from the pages in BufferPool by calling the method from HeapFile to update the page, then BufferPool updates the records by re-inserting the pages into the BufferPool. When the BufferPool is full and a new page need adding, writePage from HeapFile will be called to write the dirty page to disk and add the new page into the BufferPool. BufferPool is in charge of updating the pages because Insert/Delete operators directly call insert/delete of BufferPool.

Examples:

Workflow (join(), showing next() call):



Workflow (insertTuple()):



3. Design decisions: I chose to use nested loop join. Nested-loop-join is the slowest but it works for all kinds of comparisons (Hash join only works well for equality comparison, sort-merge join only works well for inequality comparison). For eviction policy, I chose to evict the first page every time the `BufferPool` is full. This does not work well in terms of keeping most commonly used page in memory, which result in extra disk IOs, but the implementation is easy and does not require extra data structures in order to keep track of page usage.
4. Extra Unit tests: A unit test for the correctness of `HeapFile insertTuple` behavior would be great. The spec requires the update of `BufferPool` pages to happen in the `insert/delete` methods of `BufferPool` instead of in `HeapFile`. I did it the wrong way and caused the `handleManyDirty` test to fail, and it took me a while to realize that test uses a overridden `insert` method and `BufferPool` should be in charge of updating the record.
5. Changes to API: I did not make changes to the APIs given.
6. Missing elements: I believe I finished the entire Lab 2